

Self Reflection

Roshan Shrestha

2023-02-27

Introduction

My growth in this course has always been progressive. I enjoyed the Modeling and Regression course as it helped me refresh my knowledge and concept with the statistical concepts. Through this course, I believe I have built a strong foundation with statistical analysis and statistical modeling. Furthermore, I got the opportunity to implement the models and concept in R. This is my fourth semester using R and my coding journey with R has been progressively improved as well and I feel I'm pretty much proficient in using R now.

I started with doing the import of data, pre-processing my dataset and running some exploratory data analysis for the dataset in R. I have also performed some exploratory analysis over here as many of my classmates also wanted to see which skills were frequently preferred, optimal days of job postings, the top skill set for the data analysis role, and the locations of the job posting. I have also included my visualization in my final project so that everyone can benefit from the analysis I have made on the job postings.

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
jobs <- read.csv("gsearch_jobs1.csv", row.names = 1)
```

I further performed some pre-processing on my dataset to get the desired dataset to run my model.

```
# replace string values in location column  
jobs$location <- gsub("\\+.*", "", jobs$location)
```

```
library(qdap)
```

```
## Loading required package: qdapDictionaries  
  
## Loading required package: qdapRegex
```

```
##
## Attaching package: 'qdapRegex'

## The following object is masked from 'package:dplyr':
##
##     explain

## Loading required package: qdapTools

##
## Attaching package: 'qdapTools'

## The following object is masked from 'package:dplyr':
##
##     id

## Loading required package: RColorBrewer

##
## Attaching package: 'qdap'

## The following objects are masked from 'package:base':
##
##     Filter, proportions
```

```
skill_frequecny <- freq_terms(jobs$description_tokens)
```

```
skill_frequecny
```

```
##      WORD      FREQ
## 1  'sql'      6217
## 2  'excel'    4250
## 3  'tableau'  3404
## 4  'python'   3319
## 5  'powerbi'  3186
## 6  'r'        2209
## 7  'sas'      1233
## 8  'powerpoint' 946
## 9  'word'     917
## 10 'snowflake' 845
## 11 'aws'      734
## 12 'azure'    674
## 13 'spss'     520
## 14 'jira'     499
## 15 'looker'   484
## 16 'go'       477
## 17 'microstrategy' 457
## 18 'spark'    449
## 19 'c'        437
## 20 'spreadsheet' 413
```

```
library(plotly)
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:qdapRegex':
```

```
##
```

```
##      %>%
```

```
##
```

```
## Attaching package: 'plotly'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      last_plot
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##      filter
```

```
## The following object is masked from 'package:graphics':
```

```
##
```

```
##      layout
```

```
analysis_skills <- data.frame("Categorie" = rownames(skill_frequecny), skill_frequecny)
data <- analysis_skills[, c('WORD', 'FREQ')]
```

```
colors <- c('rgb(211,94,96)', 'rgb(128,133,133)', 'rgb(144,103,167)', 'rgb(171,104,87)', 'rgb(114,147,201)')
```

```
fig <- plot_ly(data, labels = ~WORD, values = ~FREQ, type = 'pie',
  textposition = 'inside',
  textinfo = 'label+percent',
  insidetextfont = list(color = '#FFFFFF'),
  hoverinfo = 'text',
  text = ~paste(FREQ),
  marker = list(colors = colors,
    line = list(color = '#FFFFFF', width = 1)),
    #The 'pull' attribute can also be used to create space between the sectors
  showlegend = FALSE)
```

```
fig <- fig %>% layout(title = 'Pie Chart of on demand Data Analysis skills',
  xaxis = list(showgrid = FALSE, zeroline = FALSE, showticklabels = FALSE),
  yaxis = list(showgrid = FALSE, zeroline = FALSE, showticklabels = FALSE))
```

```
fig
```

```
## PhantomJS not found. You can install it with webshot::install_phantomjs(). If it is installed, please
```

```

library(plotly)
set.seed(1)

Skills <- data.frame("Categorie" = rownames(skill_frequecny), skill_frequecny)
data <- Skills[, c('WORD', 'FREQ')]

# Generate 20 random colors
colors <- paste0("rgb(", round(runif(20, 0, 255)), ",",
                  round(runif(20, 0, 255)), ",",
                  round(runif(20, 0, 255)), ")")

fig <- plot_ly(data, x = ~WORD, y = ~FREQ, type = 'bar',
               marker = list(color = colors))

fig <- fig %>% layout(title = 'Top 20 Skills set in Data Analysis Jobs',
                    xaxis = list(title = 'Skills '),
                    yaxis = list(title = 'Frequency'))

fig

```

```

job_python <- read.csv("job_python.csv", row.names = 1)
cleanjobs_salary <- filter(job_python, job_python$salary_standardized != 'NA')

```

```

library(ggplot2)
library(plotly)
library(lubridate)

```

```

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

```

```

# Convert date column to date object
job_python$date_time <- ymd_hms(job_python$date_time)

```

```

# Create a new column with the day of the week
job_python$day <- weekdays(job_python$date_time)

```

```

# Create a new column for day of the week using base R weekdays() function
job_python$day <- weekdays(job_python$date_time)

```

```

# Aggregate job counts by day of the week
job_counts <- aggregate(title ~ day, data = job_python, FUN = length)

```

```

# Create the plot using ggplot2
job_counts_plot <- ggplot(data = job_counts, aes(x = day, y = title, group = 1)) +
  geom_line(color = "steelblue", size = 1.2) +
  labs(title = "Job Posts by Day of the Week", x = "Day of the Week", y = "Number of Job Posts") +
  scale_x_discrete(limits = c("Sunday", "Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday")) +
  theme_minimal()

```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.  
## i Please use 'linewidth' instead.
```

```
# Create an interactive plot using plotly  
ggplotly(job_counts_plot)
```

```
library(tidyr)  
Jobs_lon_lat <- read.csv("Jobs_lon_lat.csv", row.names = 1)  
# Create a new data frame with the separated columns  
Jobs_lon_lat_new <- Jobs_lon_lat %>%  
  separate(location, into = c("city", "state"), sep = ", ")
```

```
## Warning: Expected 2 pieces. Missing pieces filled with 'NA' in 128 rows [3, 59,  
## 63, 167, 188, 190, 195, 228, 314, 380, 433, 482, 483, 484, 485, 507, 508, 608,  
## 617, 618, ...].
```

```
library(dplyr)  
  
# group by location and count frequency  
Jobs_lon_lat_freq <- Jobs_lon_lat %>%  
  group_by(location) %>%  
  summarise(freq = n())  
  
# join the frequency count to Jobs_lon_lat  
Jobs_lon_lat_with_freq <- left_join(Jobs_lon_lat, Jobs_lon_lat_freq, by = "location")
```

```
library(ggplot2)  
library(maps)  
library(plotly)  
  
# Filter Jobs_lon_lat dataframe to only include locations within the USA  
Jobs_lon_lat_USA <- Jobs_lon_lat_with_freq[Jobs_lon_lat_with_freq$longitude > -125 &  
  
# Create a US map using ggplot2 and maps  
USA_map <- map_data("state")  
p <- ggplot() +  
  geom_polygon(data = USA_map, aes(x = long, y = lat, group = group), fill = "white", color = "black") +  
  geom_point(data = Jobs_lon_lat_USA, aes(x = longitude, y = latitude, label = location, text = paste(")",  
  coord_map() +  
  labs(title = "Job Postings by Location in USA")
```

```
## Warning in geom_point(data = Jobs_lon_lat_USA, aes(x = longitude, y =  
## latitude, : Ignoring unknown aesthetics: label and text
```

```
ggplotly(p)
```

```
# Aggregate job counts by location and sort in descending order  
job_counts <- Jobs_lon_lat_USA %>%  
  group_by(location) %>%  
  summarize(count = n()) %>%  
  arrange(desc(count)) %>%
```

```

top_n(30, count)

# Create the plot using ggplot2
job_counts_plot <- ggplot(data = job_counts, aes(x = location, y = count)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  labs(title = "Job Posts by Location", x = "Location", y = "Number of Job Posts") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1))

# Create an interactive plot using plotly
ggplotly(job_counts_plot)

```

Removing the [] from description_tokens

```

clean_salary_desp <- filter(cleanjobs_salary, cleanjobs_salary$description_tokens != '[]')

library(qdap)
skill_frequecnny <- freq_terms(cleanjobs_salary$description_tokens)

```

How I met the Course Objectives

The course objectives were fairly simple and straightforward where we had to mainly work with the statistical modeling and the statistical analysis of the data and how each statistical measure helped us in identifying the best model. Conducting and selecting the best model, I guess was one of the toughest part where we have to understand and look at each variable but I believe that integral part has built a strong foundation on me to determine and apply appropriate models in specific data context.

Objective 1: Determine and apply the appropriate generalized linear model for a specific data context

After completing the exploratory data analysis, I worked on creating a selected dataset for my model. The below code shows how I created it for my model.

```

library(dplyr)
library(caret)

```

```
## Loading required package: lattice
```

```

pred_data <- clean_salary_desp[, c("salary_standardized", "airflow", "alteryx", "apl", "asp.net", "asser
, "bigquery", "bitbucket", "c", "c..", "c.c.."
, "cobol", "cognos", "crystal", "css", "dart"
, "datarobot", "dax", "docker", "dplyr", "excel"
, "fortran", "gcp", "gdpr", "ggplot2", "git"
, "github", "gitlab", "go", "golang", "graphql"
, "groovy", "hadoop", "html", "java", "javascript"
, "jira", "jquery", "js", "julia", "jupyter"
, "keras", "linux", "linux.unix", "looker", "matlab"
, "matplotlib", "microstrategy", "mongo", "mongodb", "mssql"

```

```

, "mxnet"           , "mysql"           , "nltk"           , "no.sql"         , "node"
, "node.js"         , "nosql"           , "nuix"           , "numpy"          , "outlook"
, "pandas"          , "perl"            , "php"            , "pl.sql"         , "plotly"
, "postgres"        , "postgresql"      , "power_bi"       , "powerpoint"     , "powerpoint"
, "powershell"      , "pyspark"         , "python"         , "pytorch"        , "qlik"
, "r"               , "redis"           , "redshift"       , "rshiny"         , "ruby"
, "rust"            , "sap"             , "sas"            , "scala"          , "scikit.l
, "seaborn"         , "selenium"        , "sharepoint"     , "shell"          , "snowflak
, "solidity"        , "spark"           , "splunk"         , "spreadsheet"    , "spss"
, "sql"             , "ssis"            , "ssrs"           , "swift"          , "t.sql"
, "tableau"         , "tensorflow"      , "terminal"       , "tidyr"          , "twilio"
, "typescript"      , "unix"            , "unix.linux"     , "vb.net"         , "vba"
, "visio"           , "visual_basic"    , "vue"            , "vue.js"         , "word" )

```

As I further did some cleaning of my data because it had little information which wouldn't be useful in predicting a model. So, the below code calculates the column sums and it selects the column of pred_data except those with a sum of 0. Hence, this code removes any columns from pred_data that have sum less than 4.

```

# calculate the column sums and find the indices of the columns with a sum of 0
zero_cols <- which(colSums(pred_data) < 4)

# select all columns except those with a sum of 0
pred_data <- pred_data[, -zero_cols]

```

Logistic Regression

Logistic Regression is used to fit the sigmoid function to the data which estimates the probability of particular observation belonging to positive class which is salary_binary. Similarly, the logistic regression is trained on the training data and is used to predict the binary outcome variable for the test data.

With logistic regression, you can predict the probability of an individual having a certain skill or not, based on the other skills they possess. This can be useful for tasks such as predicting the likelihood of a candidate being a good fit for a job based on their skillset, or identifying which skills are most important for a particular role.

I have implemented the below code for the logistic regression and accuracy is calculated to see the performance of the model.

```

pred_data1 <- pred_data

# Calculate the mean salary
mean_salary <- mean(pred_data1$salary_standardized)

# Set salary to 0 or 1 based on the mean
pred_data1$salary_binary <- ifelse(pred_data1$salary_standardized < mean_salary, 0,
                                   ifelse(pred_data1$salary_standardized >= mean_salary, 1, NA))

# Split data into train and test sets
train_index <- sample(nrow(pred_data1), 0.7 * nrow(pred_data1))
train_data <- pred_data1[train_index, ]
test_data <- pred_data1[-train_index, ]

# Build logistic regression model
log_model <- glm(salary_binary ~ ., data = train_data, family = "binomial")

```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
# Make predictions on the test set
predictions <- predict(log_model, newdata = test_data, type = "response")

# Calculate accuracy
threshold <- 0.5
predicted_classes <- ifelse(predictions > threshold, 1, 0)
accuracy <- mean(predicted_classes == test_data$salary_binary)

# Print accuracy
print(paste0("Accuracy: ", round(accuracy, 2)))
```

```
## [1] "Accuracy: 0.99"
```

Linear Regression

I have also performed linear regression in my final project where the linear regression is used to model the relationship between the salary_standardized variable and the other variables. It uses the `lm` function to fit the linear regression model by finding the best-fit line that minimizes the sum of the squared difference between the observed values of the dependent variable and the predicted value of the model.

This code divides the `pred_data` dataset into a training set and a test set. It then uses the `glm()` function to create a logistic regression model and the `predict()` function to make predictions on the test set. The `predict()` function gives the projected odds of each person having a salary above or below the threshold when the `type = "response"` argument is used.

```
# Split data into training and testing sets
set.seed(123)
trainIndex <- createDataPartition(pred_data$salary_standardized, p = .8, list = FALSE, times = 1)
trainData <- pred_data[trainIndex, ]
testData <- pred_data[-trainIndex, ]

# Create linear model
linear_model <- lm(salary_standardized ~ ., data = trainData)

# Predict on test data
predictions <- predict(linear_model, newdata = testData)
```

```
# Load required libraries
library(dplyr)
library(tidyr)
library(broom)
library(stats)

# Fit linear model
linear_model <- lm(salary_standardized ~ ., data = pred_data)

# View summary of model
linm_sum <- summary(linear_model)
linm_sum
```



```
##
## Call:
## lm(formula = salary_standardized ~ ., data = pred_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -95393 -26312    129   19273 232468
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   94602.43    1133.92   83.429 < 2e-16 ***
## airflow       8276.72    10412.39    0.795 0.426747
## alteryx       2548.64     4907.43    0.519 0.603565
## atlassian    -5577.63    15119.60   -0.369 0.712231
## aws           4879.38     4215.43    1.158 0.247170
## azure         1514.62     4891.22    0.310 0.756843
## bigquery     -6952.32     7540.29   -0.922 0.356600
## c             14197.02     5998.17    2.367 0.018009 *
## c..           9373.23    12930.06    0.725 0.468566
## cognos        5839.11     8972.19    0.651 0.515231
## crystal     -31439.79    14694.56   -2.140 0.032481 *
## dax          -7559.68     7792.57   -0.970 0.332078
## docker        8000.57    14582.65    0.549 0.583301
## excel       -10846.22     1909.77   -5.679 1.50e-08 ***
## gcp          -15121.24    13005.55   -1.163 0.245065
## ggplot2       1081.04    19221.93    0.056 0.955155
## git          13252.28    10796.73    1.227 0.219767
## github       -7688.03    12327.39   -0.624 0.532908
## go           224.69     4843.75    0.046 0.963005
## hadoop        10751.41     7062.90    1.522 0.128068
## html         -3030.75    12586.88   -0.241 0.809739
## java         14028.41     8270.15    1.696 0.089951 .
## javascript    7310.18     7569.83    0.966 0.334283
## jira          8140.17     4892.26    1.664 0.096253 .
## julia       -27414.50    20265.53   -1.353 0.176245
## jupyter      -6611.10    11293.71   -0.585 0.558342
## keras        62794.12    28605.39    2.195 0.028236 *
## linux       -17369.00    12919.04   -1.344 0.178917
## looker       6833.91     4435.87    1.541 0.123532
## matlab      -3342.84     8998.23   -0.371 0.710295
## matplotlib   -7172.36    15640.72   -0.459 0.646581
## microstrategy -10548.66     5220.88   -2.020 0.043434 *
## mongodb       11573.96    14856.17    0.779 0.436009
## mssql        39165.54    20782.99    1.885 0.059606 .
## mysql       -14497.90     7815.79   -1.855 0.063714 .
## nosql        18171.63     8253.18    2.202 0.027767 *
## numpy        19763.22    13788.41    1.433 0.151883
## outlook     -8649.96     6582.92   -1.314 0.188959
## pandas     -11798.51    12208.80   -0.966 0.333935
## pl.sql       20593.11    10417.36    1.977 0.048166 *
## postgres    -21767.39    13687.05   -1.590 0.111871
## postgresql  -9094.87    12489.03   -0.728 0.466538
## power_bi     1919.05     2373.41    0.809 0.418839
## powerpoint    7880.50     4098.26    1.923 0.054600 .
```

```
## powershell      11031.07   15928.14    0.693 0.488651
## pyspark         -179.11   13702.94   -0.013 0.989572
## python          3263.67    2547.23    1.281 0.200214
## pytorch        -29167.46   20475.74   -1.424 0.154422
## qlik            6681.14    5506.71    1.213 0.225133
## r               2911.96    2869.82    1.015 0.310349
## redshift        17314.93    7339.42    2.359 0.018388 *
## ruby           -1981.12   19470.65   -0.102 0.918964
## sap             1649.65    5115.52    0.322 0.747114
## sas             3832.62    3410.46    1.124 0.261205
## scala           2859.38   11477.67    0.249 0.803283
## scikit.learn   -3306.89   18411.01   -0.180 0.857469
## sharepoint      2070.10    5940.84    0.348 0.727528
## shell          -5548.51   21973.52   -0.253 0.800667
## snowflake       15484.67    3549.69    4.362 1.34e-05 ***
## spark           20230.41    8685.32    2.329 0.019919 *
## splunk          21069.92   19197.84    1.098 0.272515
## spreadsheet     -4923.52    3734.69   -1.318 0.187509
## spss           -16221.04    4770.58   -3.400 0.000683 ***
## sql             4209.58    1972.10    2.135 0.032887 *
## ssis            216.97     8171.39    0.027 0.978818
## ssrs            8037.97    8790.89    0.914 0.360614
## swift          11165.39    7460.49    1.497 0.134615
## t.sql           8850.78   11520.24    0.768 0.442388
## tableau        -91.57     2296.30   -0.040 0.968195
## tensorflow     -40057.70   22060.51   -1.816 0.069512 .
## terminal       -5674.81   15558.69   -0.365 0.715337
## unix           9578.68   17349.90    0.552 0.580934
## vba            -5826.71    6683.43   -0.872 0.383389
## visio           6442.98    8026.20    0.803 0.422195
## visual_basic   -4758.97   16401.78   -0.290 0.771724
## word           -7825.10    4249.77   -1.841 0.065689 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 37540 on 2669 degrees of freedom
## Multiple R-squared:  0.08721,    Adjusted R-squared:  0.06156
## F-statistic: 3.4 on 75 and 2669 DF,  p-value: < 2.2e-16
```

Training my model with my Final Project Dataset

I have performed the below code to see how my linear model is performing and here to test my function and the model, I have only selected sql skills to see the predicted salary and we can see that the predicted salary value is \$98812.01

```
# Let's assume that the user inputs the skills as a vector of strings
input_skills <- c("sql")

# Create an empty data frame with the same columns as pred_data
input_data <- data.frame(matrix(ncol = ncol(pred_data), nrow = 1))
colnames(input_data) <- colnames(pred_data)

# Set the values of the input data frame based on the user's input skills
```

```

for (skill in input_skills) {
  input_data[[skill]] <- 1
}

input_data[is.na(input_data)] <- 0

# Convert the input data frame to the same format as train_data
input_data_scaled <- data.frame(input_data[, -1])

# colnames(input_data_scaled) <- colnames(train_data_scaled)

# Use the trained model to predict the salary
predicted_salary <- predict(linear_model, newdata = input_data_scaled)

# Print the predicted salary
cat("Predicted salary: $", round(predicted_salary, 2), "\n")

```

```
## Predicted salary: $ 98812.01
```

```

# Calculate RMSE
library(Metrics)

```

```

##
## Attaching package: 'Metrics'

## The following objects are masked from 'package:caret':
##
##   precision, recall

```

```

rmse <- rmse(testData$salary, predictions)
rmse

```

```
## [1] 36902.63
```

To calculate the accuracy of a linear regression model in R, I use the coefficient of determination (R-squared) metric. R-squared measures how well the model fits the data and ranges from 0 to 1, with higher values indicating a better fit.

```

# R-squared value
rsq <- summary(linear_model)$r.squared
print(rsq)

```

```
## [1] 0.08721407
```

```

library(Metrics)

# Predict on test data
predictions <- predict(linear_model, newdata = testData)

```

```
# Calculate MAE
mae <- MAE(testData$salary, predictions)
print(paste0("MAE: ", round(mae, 2)))
```

```
## [1] "MAE: 25785.62"
```

Linear Discriminant Analysis

I have also tried implementing the linear discriminant analysis in my final project dataset through which I obtain a summary of my model's result. Discriminant Analysis is a statistical technique which is used for predicting the categorical outcomes based on a set of the predictor variable.

```
#discriminant Analysis
library(MASS)
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:plotly':
##
##      select

## The following object is masked from 'package:dplyr':
##
##      select
```

```
# Subset the data to exclude the outcome variable
predictors <- pred_data[, -which(names(pred_data) == "job_category")]

# Fit the discriminant analysis model
discriminant <- lda(salary_standardized ~ ., data = pred_data)

summary(discriminant)
```

```
##           Length Class  Mode
## prior         417  -none-  numeric
## counts         417  -none-  numeric
## means        31275  -none-  numeric
## scaling        5625  -none-  numeric
## lev           417  -none-  character
## svd             75  -none-  numeric
## N               1  -none-  numeric
## call           3  -none-   call
## terms          3  terms   call
## xlevels         0  -none-   list
```

Objective 2: Describe probability as a foundation of statistical modeling, including inference and maximum likelihood estimation

Confidence Interval

I have implemented the confidence interval for the linear model where it calculates the confidence interval for the coefficient of the linear regression model. Through the below confidence interval, it helps in identifying the uncertainty and significance of the relationship between the predictor variable and the dependent variable in the linear regression model.

```
# Calculate confidence intervals for coefficients
ci <- confint(linear_model)
ci
```

```
##           2.5 %      97.5 %
## (Intercept)  92378.9698 96825.8874
## airflow    -12140.4502 28693.8883
## alteryx     -7074.1203 12171.3915
## atlassian   -35224.9514 24069.6972
## aws         -3386.4629 13145.2158
## azure       -8076.3429 11105.5831
## bigquery    -21737.7144  7833.0706
## c           2435.4807 25958.5536
## c..         -15980.7176 34727.1848
## cognos      -11754.0278 23432.2514
## crystal     -60253.6738 -2625.9076
## dax         -22839.7724  7720.4066
## docker      -20593.8768 36595.0090
## excel       -14591.0001 -7101.4338
## gcp         -40623.2075 10380.7249
## ggplot2     -36610.3358 38772.4231
## git         -7918.5188 34423.0839
## github      -31860.2270 16484.1638
## go          -9273.2001  9722.5806
## hadoop      -3097.9047 24600.7323
## html        -27711.7742 21650.2701
## java        -2188.1447 30244.9714
## javascript  -7533.1554 22153.5100
## jira        -1452.8431 17733.1764
## julia       -67152.2359 12323.2369
## jupyter     -28756.3987 15534.2001
## keras        6703.1468 118885.0906
## linux       -42701.3444  7963.3400
## looker      -1864.1666 15531.9951
## matlab      -20987.0404 14301.3649
## matplotlib  -37841.5208 23496.7991
## microstrategy -20786.0381  -311.2766
## mongodb     -17556.8037 40704.7308
## mssql       -1586.8448 79917.9288
## mysql       -29823.5128  827.7165
## nosql        1988.3540 34354.9061
## numpy       -7273.8296 46800.2791
## outlook     -21558.1040  4258.1781
## pandas      -35738.1814 12141.1657
## pl.sql       166.2014 41020.0101
## postgres    -48605.6997  5070.9106
## postgresql  -33584.0308 15394.2929
## power_bi    -2734.8515  6572.9481
## powerpoint   -155.5889 15916.5806
```

```
## powershell      -20201.6769  42263.8199
## pyspark          -27048.5686  26690.3504
## python           -1731.0854   8258.4215
## pytorch          -69317.3889  10982.4712
## qlik             -4116.7193  17478.9970
## r                -2715.3438   8539.2648
## redshift         2923.4131  31706.4443
## ruby            -40160.2085  36197.9694
## sap              -8381.1306  11680.4395
## sas              -2854.7800  10520.0267
## scala            -19646.6500  25365.4122
## scikit.learn     -39408.1816  32794.3997
## sharepoint       -9579.0110  13719.2112
## shell            -48635.3667  37538.3449
## snowflake        8524.2423  22445.1022
## spark            3199.7692  37261.0575
## splunk           -16574.2298  58714.0690
## spreadsheet      -12246.6939  2399.6560
## spss             -25575.4412 -6866.6367
## sql              342.5796   8076.5869
## ssis             -15805.9272  16239.8736
## ssrs             -9199.6756  25275.6217
## swift            -3463.5320  25794.3032
## t.sql            -13738.7187  31440.2773
## tableau          -4594.2659   4411.1330
## tensorflow       -83315.1097   3199.7196
## terminal         -36183.1132  24833.4863
## unix            -24441.9213  43599.2905
## vba              -18931.9331   7278.5231
## visio            -9295.2311  22181.1869
## visual_basic     -36920.4502  27402.5008
## word             -16158.2786   508.0839
```

Maximum Likelihood Estimates

Finding the parameter values that maximize the likelihood of receiving the observed data under a specific statistical model is done statistically using the maximum likelihood estimation approach. The MLE estimates for linear regression represent the coefficient values that give the observed data the highest likelihood. I tried implementing the function for the maximum likelihood coefficient to find out the coefficient from the linear regression model.

```
# Calculate maximum likelihood estimates for coefficients
lm_mle <- coef(linear_model)
lm_mle
```

```
## (Intercept)      airflow      alteryx      atlassian      aws
## 94602.42863      8276.71909      2548.63563     -5577.62708      4879.37644
##      azure      bigquery      c      c..      cognos
## 1514.62007     -6952.32189     14197.01718     9373.23361     5839.11180
##      crystal      dax      docker      excel      gcp
## -31439.79070     -7559.68288     8000.56608    -10846.21699    -15121.24130
##      ggplot2      git      github      go      hadoop
## 1081.04367     13252.28259     -7688.03162     224.69025     10751.41382
##      html      java      javascript      jira      julia
```

```
## -3030.75207 14028.41332 7310.17731 8140.16662 -27414.49947
## jupyter keras linux looker matlab
## -6611.09928 62794.11871 -17369.00220 6833.91424 -3342.83775
## matplotlib microstrategy mongodb mssql mysql
## -7172.36082 -10548.65737 11573.96356 39165.54197 -14497.89811
## nosql numpy outlook pandas pl.sql
## 18171.63003 19763.22475 -8649.96292 -11798.50781 20593.10573
## postgres postgresql power_bi powerpoint powershell
## -21767.39456 -9094.86895 1919.04833 7880.49585 11031.07151
## pyspark python pytorch qlik r
## -179.10910 3263.66808 -29167.45887 6681.13887 2911.96054
## redshift ruby sap sas scala
## 17314.92870 -1981.11958 1649.65444 3832.62331 2859.38106
## scikit.learn sharepoint shell snowflake spark
## -3306.89093 2070.10012 -5548.51090 15484.67226 20230.41336
## splunk spreadsheet spss sql ssis
## 21069.91957 -4923.51897 -16221.03892 4209.58326 216.97320
## ssrs swift t.sql tableau tensorflow
## 8037.97303 11165.38559 8850.77933 -91.56644 -40057.69502
## terminal unix vba visio visual_basic
## -5674.81345 9578.68461 -5826.70501 6442.97787 -4758.97473
## word
## -7825.09734
```

Objective 3: Conduct model selection for a set of candidate models.

To look for the model selection, we could refer to various **goodness-of-fit** measures such as R-square, adjusted R-square, Akaike Information Criterion (AIC) or Bayesian Information Criterion (BIC). I implemented the statistical modeling through the below code which provides various statistics for each coefficient. The tidy format helps in easily manipulating, visualizing, and analyzing the model output. Furthermore, with the values we get from the *AIC* and *BIC*, it helps in assessing and comparing different models based on their goodness of fit and complexity.

```
library(broom)
```

```
tidy(linear_model)
```

```
## # A tibble: 76 x 5
##   term      estimate std.error statistic p.value
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept) 94602.    1134.    83.4      0
## 2 airflow     8277.    10412.    0.795    0.427
## 3 alteryx     2549.    4907.    0.519    0.604
## 4 atlassian  -5578.    15120.   -0.369    0.712
## 5 aws        4879.    4215.    1.16     0.247
## 6 azure      1515.    4891.    0.310    0.757
## 7 bigquery   -6952.    7540.   -0.922    0.357
## 8 c         14197.    5998.    2.37     0.0180
## 9 c..       9373.    12930.    0.725    0.469
## 10 cognos     5839.    8972.    0.651    0.515
## # ... with 66 more rows
```

```
glance(linear_model)
```

```
## # A tibble: 1 x 12
##   r.squared adj.r.^1 sigma stati~2 p.value    df logLik    AIC    BIC devia~3
##   <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>    <dbl> <dbl> <dbl>    <dbl>
## 1    0.0872  0.0616 37538.    3.40 3.39e-20    75 -32770. 65694. 66149. 3.76e12
## # ... with 2 more variables: df.residual <int>, nobs <int>, and abbreviated
## #   variable names 1: adj.r.squared, 2: statistic, 3: deviance
```

Cross-Validation

Cross-validation is used to for estimating the performance of the predictive model. Cross-validation involves dividing the data into two parts. The validation set is used to estimate the model's performance on new or the unseen data. We mainly use cross-validation to avoid the overfitting as implementing the cross-validation helps in getting more accurate estimate of the model's performance.

Here, I have implemented the cross validation where I have set up a 10-fold cross validation which means that the data will be split into 10 equal parts and the model will be trained on 9 parts and tested on the remaining part.

```
# training a linear regression model using 10-fold cross-validation
control <- trainControl(method = "cv", number = 10)
Fit <- train(salary_standardized ~ ., data = trainData, method = "lm", metric = "RMSE", trControl = control)
# pred_data$salary_standardized

print(Fit)
```

```
## Linear Regression
##
## 2198 samples
## 75 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1979, 1979, 1979, 1977, 1979, 1978, ...
## Resampling results:
##
##   RMSE      Rsquared    MAE
## 38922.76  0.04047155 28356.68
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

From the above output, we can compute that the EMSE is the measure of the average difference between the predicted value and the actual value. Similarly, by looking at the R-squared value, mean absolute error, RMSE we can have a model selection based on that. The R-squared measures the proportion of the variance while the mean absolute error measure the absolute difference between the predicted and actual values.

Calculating RMSE

This code calculates the root mean square between the predicted salaries and the actual salaries in the rest dataset. The difference between the predicted and actual salaries is first squared, then the mean of these squared differences is taken, and finally the square root of this mean is calculated.


```

# Calculate RMSE
mse <- mean((predictions - testData$salary_standardized)^2)
rmse <- sqrt(mse)
rmse

```

```
## [1] 34833.4
```

Bootstrapping

```

library(caret)

# Identify near zero variance predictors in the data frame
nzv <- nearZeroVar(pred_data, saveMetrics = TRUE)

# Remove the identified near zero variance predictors
pred_data <- pred_data[, !nzv$nzv]

lda_model <- lda(salary_standardized ~ ., data = pred_data)

```

```

#Bootstrapping
library(MASS)
library(boot)

```

```

##
## Attaching package: 'boot'

## The following object is masked from 'package:lattice':
##
##      melanoma

```

```

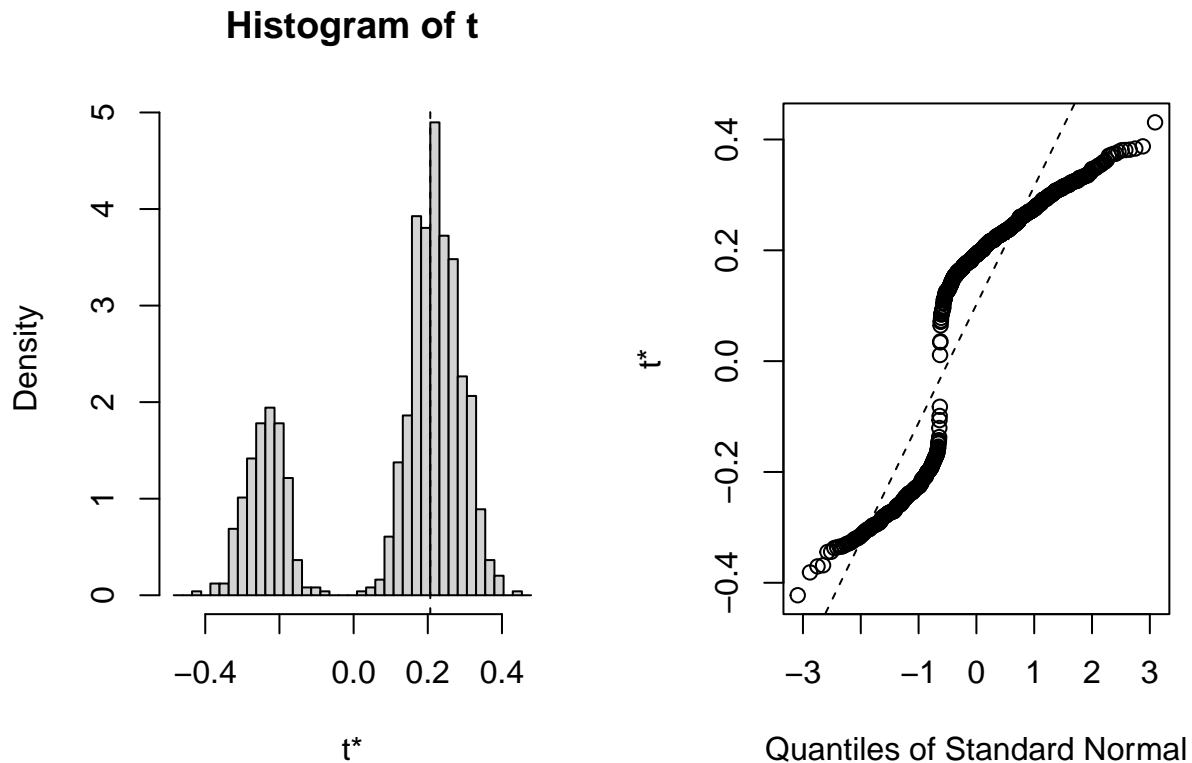
# Subset the data to exclude the outcome variable
predictors <- pred_data[, -which(names(pred_data) == "job_category")]

lda_coef <- function(data, index) {
  fit <- lda(salary_standardized ~ ., data = data[index, ])
  coef(fit)
}

# Use bootstrapping to estimate the standard errors of the coefficients
set.seed(123) # for reproducibility
boot_lda <- boot(data = pred_data, statistic = lda_coef, R = 1000)

plot(boot_lda)

```




Objective 4: Communicate the results of statistical models to a general audience

Communicating the result plays an important role and when communicating the result, it should be very precise so that the audience can easily understand it. It is better to use plain language and while there might be some of the technical concepts and technical terminology, it is better to provide context to the audience. I'm a great believer of visual aids such as graph, chart, table which could easily grab an attention of the audience and they can easily understand the findings.

The below plot is one of the visual aid for my linear model. It is a residual plot which is used to check the assumption of the linear model. The horizontal axis represents the predicted values and the vertical axis represents the residuals or the difference between the observed values and the predicted values. As this plot shows a random scatter points we can compute that my linear model met the assumption.

```
R_fit <- augment(linear_model)

# plot fitted values and residuals
ggplot(data = R_fit, aes(x = .fitted, y = .resid)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  xlab("Fitted values") +
  ylab("Residuals")
```




The below code is another method of communicating the result to the audience which is through the table data. This helps us in showing the importance of each variable as per the ranking.

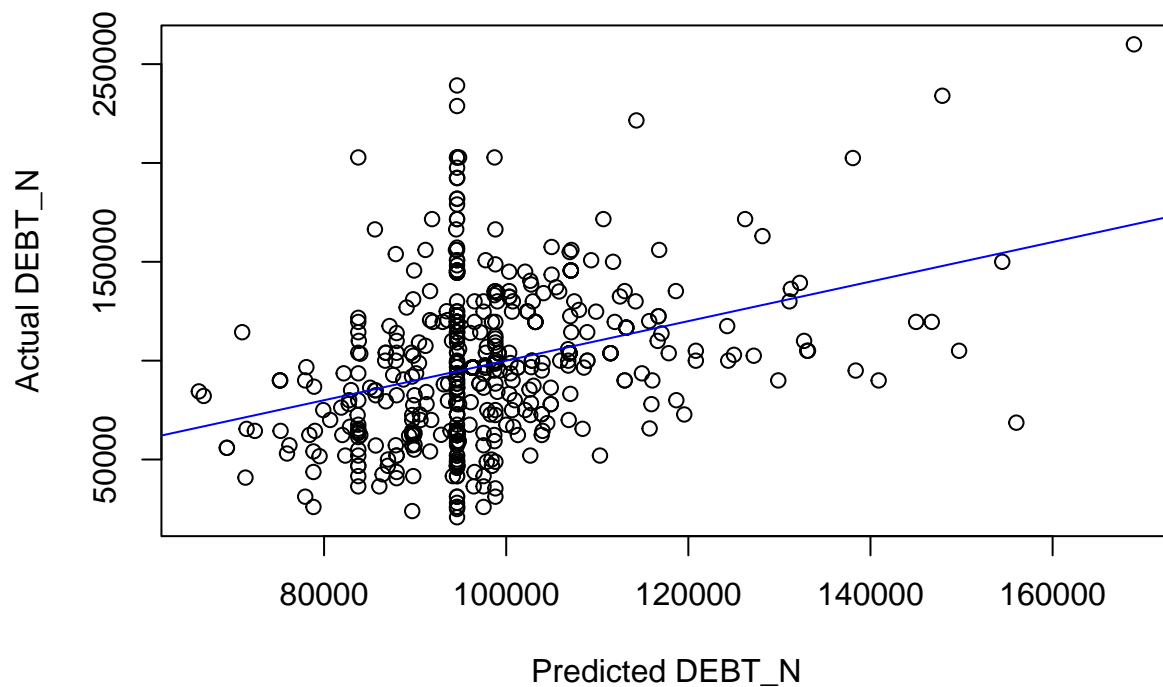
```
Imp <- varImp(Fit, scale = FALSE)
print(Imp)
```

```
## lm variable importance
##
##   only 20 most important variables shown (out of 75)
##
##           Overall
## excel          4.653
## snowflake      3.239
## spss           2.926
## c              2.730
## spark          2.537
## redshift       2.201
## docker         2.081
## mssql          1.953
## java           1.950
## crystal        1.920
## lookout        1.808
## microstrategy  1.716
## airflow        1.714
## jira           1.701
## nosql          1.696
```

```
plot(Imp)
```




20



I have also used a QQ plot in my final project to have a overview of how my model is performing and to check the assumptions for my linear model. QQ plot helps in visually comparing the distribution of the dataset which is also useful for identifying deviations from normality, detecting outliers, and assess the overall shape of the data distribution.

```
# Get the residuals
residuals <- residuals(linear_model)

# QQ plot
qqnorm(residuals)
qqline(residuals)
```



Objective 5: Use programming software (i.e. R) to fit and assess statistical models

Apart from implementing linear regression, logistic regression and discriminant analysis, I also tried implementing other various machine learning algorithm in my dataset to see how each different machine learning models performed.

Random Forest

I implemented the below code to see how it predicts for the regression analysis. I have used the `randomForest` function to fit the random forest model and it then calculates the mean absolute error to evaluate the model's performance.

```
library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
## The following object is masked from 'package:dplyr':  
##  
##      combine
```

```
set.seed(123)  
# Split data into train and test sets  
train_index <- sample(nrow(pred_data), 0.7 * nrow(pred_data))  
train_data <- pred_data[train_index, ]  
test_data <- pred_data[-train_index, ]  
  
# Fit random forest model  
rf_model <- randomForest(salary_standardized ~ ., data = train_data, ntree = 75, mtry = sqrt(ncol(train_data)))  
  
# Make predictions on the test data  
predictions <- predict(rf_model, test_data)  
  
# Calculate the mean absolute error  
mae <- mean(abs(predictions - test_data$salary_standardized))  
print(paste0("MAE: ", round(mae, 2)))
```

```
## [1] "MAE: 27567.19"
```

XGBOOST

I also tried implementing the XGBoost and tried to see the same result on how the model performed through fitting the model and finding the mean absolute error.

```
library(xgboost)
```

```
##  
## Attaching package: 'xgboost'
```

```
## The following object is masked from 'package:plotly':  
##  
##      slice
```

```
## The following object is masked from 'package:dplyr':  
##  
##      slice
```

```
# Split data into train and test sets  
train_index <- sample(nrow(pred_data), 0.7 * nrow(pred_data))  
train_data <- pred_data[train_index, ]  
test_data <- pred_data[-train_index, ]  
  
# Set up xgboost matrix  
xg_train <- xgb.DMatrix(as.matrix(train_data[, -1]), label = train_data$salary_standardized)  
xg_test <- xgb.DMatrix(as.matrix(test_data[, -1]), label = test_data$salary_standardized)  
  
# Define hyperparameters  
params <- list(  
  objective = "reg:squarederror",
```

```

eta = 0.1,
max_depth = 6,
subsample = 0.8,
colsample_bytree = 0.8,
min_child_weight = 3,
nthread = 4
)

# Train the model
xgb_model <- xgb.train(
  params = params,
  data = xg_train,
  nrounds = 1000,
  watchlist = list(train = xg_train, test = xg_test),
  early_stopping_rounds = 10,
  verbose = 1
)

```

```

## [1] train-rmse:96020.964474 test-rmse:93977.405804
## Multiple eval metrics are present. Will use test_rmse for early stopping.
## Will train until test_rmse hasn't improved in 10 rounds.
##
## [2] train-rmse:88097.680713 test-rmse:86009.915517
## [3] train-rmse:81074.660717 test-rmse:78992.486075
## [4] train-rmse:74819.956821 test-rmse:72726.828975
## [5] train-rmse:69376.697053 test-rmse:67243.279901
## [6] train-rmse:64690.134319 test-rmse:62509.749484
## [7] train-rmse:60508.315487 test-rmse:58315.733676
## [8] train-rmse:56950.965177 test-rmse:54794.452672
## [9] train-rmse:53890.448203 test-rmse:51759.291708
## [10] train-rmse:51273.947813 test-rmse:49141.791289
## [11] train-rmse:48995.768933 test-rmse:46910.634849
## [12] train-rmse:47128.160359 test-rmse:45072.054185
## [13] train-rmse:45475.906196 test-rmse:43457.174830
## [14] train-rmse:44115.155108 test-rmse:42167.571793
## [15] train-rmse:42952.980757 test-rmse:41043.255626
## [16] train-rmse:42003.083490 test-rmse:40155.308343
## [17] train-rmse:41209.517964 test-rmse:39416.541371
## [18] train-rmse:40562.671356 test-rmse:38811.084027
## [19] train-rmse:40019.672351 test-rmse:38324.546546
## [20] train-rmse:39555.123575 test-rmse:37926.341045
## [21] train-rmse:39142.506116 test-rmse:37567.713271
## [22] train-rmse:38814.647607 test-rmse:37274.830352
## [23] train-rmse:38570.967583 test-rmse:37079.322012
## [24] train-rmse:38348.979761 test-rmse:36906.736672
## [25] train-rmse:38161.833708 test-rmse:36766.393295
## [26] train-rmse:38021.502608 test-rmse:36660.898206
## [27] train-rmse:37890.194362 test-rmse:36556.609104
## [28] train-rmse:37799.752530 test-rmse:36507.090796
## [29] train-rmse:37708.812529 test-rmse:36432.474030
## [30] train-rmse:37616.912460 test-rmse:36398.688892
## [31] train-rmse:37546.698973 test-rmse:36375.187311
## [32] train-rmse:37486.279222 test-rmse:36347.119599

```



```
## [33] train-rmse:37451.311796 test-rmse:36333.674719
## [34] train-rmse:37413.284153 test-rmse:36298.362195
## [35] train-rmse:37379.921109 test-rmse:36279.097538
## [36] train-rmse:37348.650128 test-rmse:36270.241732
## [37] train-rmse:37321.478766 test-rmse:36274.198836
## [38] train-rmse:37304.975568 test-rmse:36268.267291
## [39] train-rmse:37288.446425 test-rmse:36260.234863
## [40] train-rmse:37268.247789 test-rmse:36267.342676
## [41] train-rmse:37253.598762 test-rmse:36273.011713
## [42] train-rmse:37241.606795 test-rmse:36269.352969
## [43] train-rmse:37230.073626 test-rmse:36272.121434
## [44] train-rmse:37219.686886 test-rmse:36273.704319
## [45] train-rmse:37213.842137 test-rmse:36266.278565
## [46] train-rmse:37203.748594 test-rmse:36269.870040
## [47] train-rmse:37195.645609 test-rmse:36271.315259
## [48] train-rmse:37189.421663 test-rmse:36277.955429
## [49] train-rmse:37183.046888 test-rmse:36284.884097
## Stopping. Best iteration:
## [39] train-rmse:37288.446425 test-rmse:36260.234863
```

```
# Make predictions on the test data
predictions <- predict(xgb_model, xg_test)

# Calculate the mean absolute error
mae <- mean(abs(predictions - test_data$salary_standardized))
print(paste0("MAE: ", round(mae, 2)))
```

```
## [1] "MAE: 26650.92"
```

I also tried implementing my cod of the XGBoost to predict the salary of the skills for the sql which showed a slightly different salary than that of the linear model.

```
# Let's assume that the user inputs the skills as a vector of strings
input_skills <- c("sql")

# Create an empty data frame with the same columns as pred_data
input_data <- data.frame(matrix(ncol = ncol(pred_data), nrow = 1))
colnames(input_data) <- colnames(pred_data)

# Set the values of the input data frame based on the user's input skills
for (skill in input_skills) {
  input_data[[skill]] <- 1
}

input_data[is.na(input_data)] <- 0

# Convert the input data frame to the same format as train_data
input_data_scaled <- data.frame(input_data[, -1])

# Predict salary using xgboost model
```

```
input_xg <- xgb.DMatrix(as.matrix(input_data_scaled))
predicted_salary <- predict(xgb_model, input_xg)

# Print predicted salaries
print(paste0("Predicted Salary: ", round(predicted_salary, 2)))
```

```
## [1] "Predicted Salary: 96306.84"
```

Support vector machine

```
library(e1071)

# Split data into train and test sets
train_index <- sample(nrow(pred_data), 0.7 * nrow(pred_data))
train_data <- pred_data[train_index, ]
test_data <- pred_data[-train_index, ]

# Train SVM model
svm_model <- svm(salary_standardized ~ ., data = train_data)

# Make predictions on the test data
predictions <- predict(svm_model, test_data)

# Calculate the mean absolute error
mae <- mean(abs(predictions - test_data$salary_standardized))
print(paste0("MAE: ", round(mae, 2)))
```

```
## [1] "MAE: 26891.75"
```

Support Vector Regression (SVR)

```
library(e1071)

# Split data into train and test sets
train_index <- sample(nrow(pred_data), 0.7 * nrow(pred_data))
train_data <- pred_data[train_index, ]
test_data <- pred_data[-train_index, ]

# Define the formula for the SVR
formula <- as.formula("salary_standardized ~ .")

# Train the SVR model
svm_model <- svm(formula, data=train_data, kernel="radial")

# Make predictions on the test data
predictions <- predict(svm_model, test_data)

# Calculate the mean absolute error
mae <- mean(abs(predictions - test_data$salary_standardized))
print(paste0("MAE: ", round(mae, 2)))
```

```
## [1] "MAE: 27043.18"
```

Among all the machine learning models I tried Support Vector machine had the lower mean absolute error.

Conclusion

I believe I have met all the objective of this course as this course has been a great learning experience for me to hone my skills and brush my theroteical concepts regarding the statistical modeling. I have attached my codes and the description on hiw I have met all the objective of the course above.